

Individual_Project

December 1, 2019

0.1 Goal:

1. This project is going to predict the number of points based on wine reviews:

$f: X \rightarrow y$

- X will be the input reviews
- y will be value of points

0.2 ##### 2. And construct a small recommendation system to recommend 5 wineries for a customer.

0.3 Roadmap:

1. **Import data:** this project is done on Google Colab, I save the raw dataset downloaded from Kaggle website on Google Drive and then import it to Google Colab.

2. Data exploration:

- Overall dataset analysis: check number of missing values, data type of each features;
- Numerical features: descriptive stats, distribution, correlation analysis;
- Categorical features: unique value, distribution analysis.

3. Data preprocessing:

- Missing data imputation;
- Drop columns;
- Categorical features: convert long sentences feature to its length, One-hot encoding;
- Reduce dimensionality: PCA.

4. Model training and validation:

- KNN: 5-fold cross validation, tuning parameter K;
- CatBoost: feature importance discussion, reduce data dimensionality by dropping unimportant features;

5. Small recommendation system – collaborative filtering:

-

0.4 Cosine similarity calculation: KNN.

0.5 Results

- KNN model regression analysis: when $k = 30$, the RMSE of testing data is 2.112;
- CatBoost model: the RMSE of testing data is 1.844, and after dropping unimportant features, the RMSE of testing data is 1.868 (doesn't change much);
- KNN recommendation system: this model only works on customers who have already reviewed on any wines, but doesn't work on new customers.

1 Part 0: Setup Google Drive Environment

```
[0]: !pip install -U -q PyDrive

from pydrive.auth import GoogleAuth
from pydrive.drive import GoogleDrive
from google.colab import auth
from oauth2client.client import GoogleCredentials

auth.authenticate_user()
gauth = GoogleAuth()
gauth.credentials = GoogleCredentials.get_application_default()
drive = GoogleDrive(gauth)
```

```
[0]: link = 'https://drive.google.com/open?id=1JU1ROzMJktzwvgdbajWOHZMEFzKi8ZDI'
fluff, id = link.split('=')
file = drive.CreateFile({'id':id})
file.GetContentFile('winemag-data_first150k.csv')
```

2 Part 1: Data Exploartion

2.0.1 Part 1.1: Understand the Raw Dataset

```
[0]: import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import numpy as np
import imblearn
```

```
# will show all the columns
pd.set_option('display.max_columns', None)

wine_df = pd.read_csv('winemag-data_first150k.csv')
```

```
[0]: wine_df = wine_df.drop('Unnamed: 0', axis = 1)
wine_df.head()
```

```
[0]:
```

	country	description \
0	Italy	Aromas include tropical fruit, broom, brimston...
1	Portugal	This is ripe and fruity, a wine that is smooth...
2	US	Tart and snappy, the flavors of lime flesh and...
3	US	Pineapple rind, lemon pith and orange blossom ...
4	US	Much like the regular bottling from 2012, this...

	designation	points	price	province \
0	Vulkà Bianco	87	NaN	Sicily & Sardinia
1	Avidagos	87	15.0	Douro
2	NaN	87	14.0	Oregon
3	Reserve Late Harvest	87	13.0	Michigan
4	Vintner's Reserve Wild Child Block	87	65.0	Oregon

	region_1	region_2	taster_name \
0	Etna	NaN	Kerin O'Keefe
1	NaN	NaN	Roger Voss
2	Willamette Valley	Willamette Valley	Paul Gregutt
3	Lake Michigan Shore	NaN	Alexander Peartree
4	Willamette Valley	Willamette Valley	Paul Gregutt

	taster_twitter_handle	title \
0	@kerinokeefe	Nicosia 2013 Vulkà Bianco (Etna)
1	@vossroger	Quinta dos Avidagos 2011 Avidagos Red (Douro)
2	@paulgwine	Rainstorm 2013 Pinot Gris (Willamette Valley)
3	NaN	St. Julian 2013 Reserve Late Harvest Riesling ...
4	@paulgwine	Sweet Cheeks 2012 Vintner's Reserve Wild Child...

	variety	winery
0	White Blend	Nicosia
1	Portuguese Red	Quinta dos Avidagos
2	Pinot Gris	Rainstorm
3	Riesling	St. Julian
4	Pinot Noir	Sweet Cheeks

```
[0]: print ("Num of rows: " + str(wine_df.shape[0])) # row count
print ("Num of columns: " + str(wine_df.shape[1])) # col count
```

Num of rows: 129971

Num of columns: 13


```
[0]: wine_df.dtypes
```

```
[0]: country          object
description          object
designation          object
points              int64
price              float64
province            object
region_1            object
region_2            object
taster_name         object
taster_twitter_handle object
title               object
variety             object
winery              object
dtype: object
```

Some columns contain missing values, let's see the number of non-missing values in each column:

```
[0]: # Some columns contain missing values, let's see the number of non-missing
      ↪ values in each column:
wine_df.describe(include='all',).T
import missingno as msno
import seaborn as sns
msno.bar(wine_df,color= sns.color_palette('viridis'))# missing values
```

```
[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f22b6fdd0b8>
```



2.1 Part 1.2 Understand the Features

2.1.1 Part 1.2.1 Descriptive Statistics of Numeric Features

```
[0]: wine_df.describe()
```


```
[0]:
```

	points	price
count	129971.000000	120975.000000
mean	88.447138	35.363389
std	3.039730	41.022218
min	80.000000	4.000000
25%	86.000000	17.000000
50%	88.000000	25.000000
75%	91.000000	42.000000
max	100.000000	3300.000000

```
[0]: %matplotlib inline
import matplotlib.pyplot as plt
import seaborn as sns

# Check distribution of 'points':
sns.distplot(wine_df['points'], bins = 20)
```


```
[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f22ba3445c0>
```




The range of points is in [80, 100], most of the wines have points less than 90.

```
[0]: # Distribution of 'price':  
sns.distplot(wine_df['price'].dropna())
```

```
[0]: <matplotlib.axes._subplots.AxesSubplot at 0x7f22b6ea6d68>
```




```
[0]: # Distribution of log transformed 'price'  
ax = sns.distplot(np.log(wine_df['price']).dropna())  
ax.set(xlabel = 'log(price)')  
plt.show()
```



The range of 'price' is in [4, 3300], however, only a small number of wines have price higher than 45. After log trasformation, the distribution seems closer to normal distribution.

```
[0]: # Correlation between 'points' and 'log(price)':  
df_plot = pd.DataFrame(data={'points': wine_df['points'], 'log(price)': np.  
    ↳log(wine_df['price'])})  
sns.jointplot(x = 'points', y = 'log(price)', data = df_plot, kind = 'reg');
```



Generally speaking, expensive wines tend to have higher points.

2.1.2 Part 1.2.2 Distribution of Categorical Features

1. Check number of unique values in each column

```
[0]: wine_df.nunique()
```

```
[0]: country          43
description        119955
designation         37979
points              21
```



```

price                390
province             425
region_1            1229
region_2              17
taster_name          19
taster_twitter_handle 15
title               11840
variety              707
winery              16757
dtype: int64

```

‘description’ and ‘title’ columns have more than 91% unique values which are not informative. Thus I remove them in data preprocessing part later.

2. Due to limited space, only some of the categorical variables are displayed here. I use barplot to display the top 20 frequent values for each variable.

```


[0]: plt.bar(x = wine_df['country'].value_counts()[0:19].index, height = wine_df['country'].value_counts()[0:19].values)
plt.xticks(rotation=90)
plt.xlabel('country')

```

```

[0]: Text(0.5, 0, 'country')


```



US, France and Italy are the three main countries where the wines come from.

```
[0]: plt.bar(x = wine_df['variety'].value_counts()[0:19].index, height =  
    ↪ wine_df['variety'].value_counts()[0:19].values)  
plt.xticks(rotation=90)  
plt.xlabel('variety')
```

```
[0]: Text(0.5, 0, 'variety')
```




Pinot Noir is the most popular variety, and chrdonnay is the second popular one.

```
[0]: plt.bar(x = wine_df['winery'].value_counts()[0:19].index, height =  
    ↪ wine_df['winery'].value_counts()[0:19].values)  
plt.xticks(rotation=90)
```

```
plt.xlabel('winery')
```

```
[0]: Text(0.5, 0, 'winery')
```




There are many unique values in column 'winery', most of them occur less than 200 times. Thus this is not a very informative column.

2.1.3 Part 1.2.3 Descriptive Statistics of a Special Column: 'description'

I tend to convert description column to its length (a new numeric column) in Part 1.2.3, before doing that, let's analyze if there is any correlation between length of descriptions and points.

```
[0]: description_to_plot = wine_df['description'].apply(lambda x: len(x.split()))
description_plot = pd.DataFrame(data={'points': wine_df['points'], 'length of_
↳description': description_to_plot})
```

```
sns.jointplot(x = 'points', y = 'length of description', data = data,
              description_plot, kind = 'reg');
```



3 Part 2 Data Preprocessing

I preprocess **2** groups of data ($[X, y]$ and $[X_2, y_2]$): 1. The first group of data $[X, y]$ is for **KNN** model training. Generally speaking, I perform One-Hot Encoding on categorical variables and then apply Principle Component Analysis on that high dimensional data. 2. Another group of data $[X_2, y_2]$ is for the second model. Because this dataset contains a lot of categorical variables, the second model I choose is **Catboost**, which is a utility model relates to a gradient lifting algorithm library, and can deal with class type features well.

3.1 Part 2.1 Preprocess Data for KNN Model

3.1.1 Part 2.1.1 Drop Columns

Remove the first index column and columns 'designation', 'region_2', 'taster_name', 'taster_twitter_handle' which have more than 20% missing values. Also, remove 'title' and 'description' columns which have more than 91% unique values.

```
[0]: wine_prep = wine_df
to_drop = ['designation', 'region_2', 'taster_name', 'taster_twitter_handle',
          ↪ 'title', 'description']
wine_prep = wine_prep.drop(to_drop, axis=1)
```

3.1.2 Part 2.1.2 Preprocess Numeric Features

1. replacing missing values in column 'points' and 'price'.

```
[0]: wine_prep['points'] = (wine_prep['points'].fillna(wine_df['points'].mean()))
          ↪ astype('int64')
wine_prep['price'] = (wine_prep['price'].fillna((wine_df['price'].mean()))).
          ↪ astype('float64')
wine_prep.head()
```

```
[0]:
```

	country	points	price	province	region_1 \
0	Italy	87	35.363389	Sicily & Sardinia	Etna
1	Portugal	87	15.000000	Douro	NaN
2	US	87	14.000000	Oregon	Willamette Valley
3	US	87	13.000000	Michigan	Lake Michigan Shore
4	US	87	65.000000	Oregon	Willamette Valley

	variety	winery
0	White Blend	Nicosia
1	Portuguese Red	Quinta dos Avidagos
2	Pinot Gris	Rainstorm
3	Riesling	St. Julian
4	Pinot Noir	Sweet Cheeks

3.1.3 Part 2.1.3 Preprocess Categorical Features

1. replacing values that don't appear frequently with 'others'
2. applying One-Hot Encoding on categorical features
3. applying PCA on features to reduce feature dimension

```
[0]: # Replace non-frequent values with 'Others':
def others(x):
    to_replace = x.value_counts().index[(x.value_counts() < 100)]
```

```

    x[x.apply(lambda x_1: x_1 in to_replace)] = 'Others'
    return x
enc_cols = wine_prep.drop(['price', 'points'], axis = 1)
enc_cols = enc_cols.apply(others)

# One-Hot encoding:
dummies = pd.get_dummies(enc_cols)

# Combined with 'price':
wine_encoded = pd.concat([wine_prep['price'], dummies], axis = 1)

```

3.1.4 Part 2.1.4 Replace Feature ‘description’ with Its Length

```

[0]: description_len = wine_df['description'].apply(lambda x: len(x.split()))
     wine_encoded = pd.concat([wine_encoded, description_len], axis = 1)

```

3.1.5 Part 2.1.5 Principal Component Analysis

Because the dimension of encoded dataset for KNN model is very high, the training speed will be slow. After applying PCA method, the final training dataset contains 50 features.

```

[0]: from sklearn import decomposition

     X = wine_encoded
     y = wine_df['points']

     # PCA
     pca = decomposition.PCA(n_components=50)
     pca.fit(X)
     X = pca.transform(X)
     print('The dimension of the first training dataset is:')
     X.shape

```

The dimension of the first training dataset is:

```

[0]: (129971, 50)

```

3.2 Part 2.2 Preprocess Data for Catboost

```

[0]: wine_prep2 = wine_df
     y2 = (wine_prep2['points'].fillna(wine_df['points'].mean())).astype('int64')
     wine_prep2['price'] = (wine_prep2['price'].fillna((wine_df['price'].mean()))).
     ↪astype('float64')
     to_drop = ['description', 'points']

```

```
wine_prep2 = wine_prep2.drop(to_drop, axis=1)
X2 = pd.concat([wine_prep2, description_len], axis = 1)
X2=X2.fillna(-1)
print('The dimension of the second training dataset is:')
X2.shape
```

The dimension of the second training dataset is:

```
[0]: (129971, 12)
```

4 Part 3: Regression Model Analysis

4.1 Part 3.1 K-Nearest Neighbors

4.1.1 Part 3.1.1 Split Dataset

```
[0]: # Split data into training and testing
from sklearn import model_selection

# Reserve 20% for testing
X_train, X_test, y_train, y_test = model_selection.train_test_split(X, y,
    ↳test_size=0.2)

print('training data has %d observation with %d features'% X_train.shape)
print('test data has %d observation with %d features'% X_test.shape)
```

training data has 103976 observation with 50 features

test data has 25995 observation with 50 features

4.1.2 Part 3.1.2 Model Training and Validation

```
[0]: # Build model
from sklearn.neighbors import KNeighborsRegressor

# K Nearest Neighbors
regressor_KNN = KNeighborsRegressor()
```

Use 5-fold Cross Validation to get the mean squared error for KNN:

```
[0]: # Use 5-fold Cross Validation to get the mean squared error for different models
model_names = ['KNN']
model_list = [regressor_KNN]
count = 0
```

```

for regressor in model_list:
    cv_score = model_selection.cross_val_score(regressor, X, y, cv=5,
    →scoring='neg_mean_squared_error')
    cv_score = np.sqrt(-cv_score)
    print(cv_score)
    print('Model RMSE of %s is: %.3f'%(model_names[count], cv_score.mean()))

```

```

[2.21111219 2.22051193 2.14503297 2.2102439 2.2033381 ]
Model RMSE of KNN is: 2.198

```

4.2 Part 3.2 Tuning Model Parameters for KNN

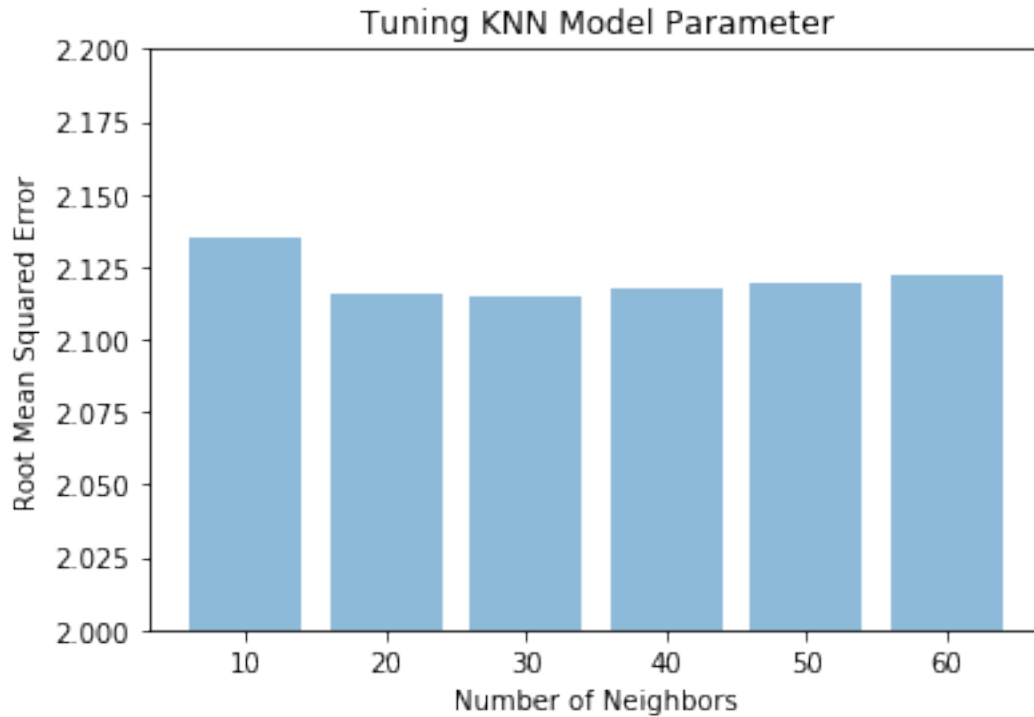
Change the number of neighbors in the KNN model (I don't use cross validation here because it takes very long time to run):

```

[0]: # change the number of neighbors (10, 20, 30, 40, 50, 60) in the KNN model
from sklearn.metrics import mean_squared_error
rmse_KNN = []
for n in [10, 20, 30, 40, 50, 60]:
    regressor_KNN = KNeighborsRegressor(n_neighbors = n)
    regressor_KNN.fit(X_train, y_train)
    y_predict_KNN = regressor_KNN.predict(X_test)
    rmse_KNN.append((mean_squared_error(y_test, y_predict_KNN))*0.5)

import matplotlib.pyplot as plt
plt.bar(np.arange(6), rmse_KNN, align='center', alpha=0.5)
plt.ylabel('Root Mean Squared Error')
plt.xlabel('Number of Neighbors')
plt.title('Tuning KNN Model Parameter')
plt.xticks(np.arange(6), [10, 20, 30, 40, 50, 60])
axes = plt.gca()
axes.set_ylim([2.0, 2.2])
plt.show()
print('When k = 30, KNN model has the best estimation with root mean squared_
    →error = %.3f' % min(rmse_KNN) )

```

When $k = 30$, KNN model has the best estimation with root mean squared error = 2.115

4.3 Part 3.3 CatBoost

Cat Boost has two advantages: 1. Firstly, it deals with categorical features in the training process rather than in the pre-processing stage. 2. Secondly, the algorithm for calculating leaf nodes can avoid overfitting when choosing tree structure.

4.3.1 Part 3.3.1 Split Dataset

```
[0]: categorical_features_indices = [0,1,3,4,5,6,7,8,9,10]

X_train2, X_test2, y_train2, y_test2 = train_test_split(X2, y2, test_size=0.2,
                                                         random_state=42)

X_train2, X_valid2, y_train2, y_valid2 = train_test_split(X_train2, y_train2,
                                                         ↪test_size=0.2,
                                                         random_state=52)
```

4.3.2 Part 3.3.2 Cat Boost Model Training and Testing

```
[0]: def perform_model(X_train, y_train, X_valid, y_valid, X_test, y_test):
    model = CatBoostRegressor(
        random_seed = 400,
        loss_function = 'RMSE',
        iterations=400,
    )

    model.fit(
        X_train, y_train,
        cat_features = categorical_features_indices,
        eval_set=(X_valid, y_valid),
        verbose=False
    )

    y_train_pred = model.predict(X_train)
    rmse_train = (np.mean((y_train - y_train_pred) **2)) **0.5

    y_test_pred = model.predict(X_test)
    rmse_test = (np.mean((y_test - y_test_pred) **2)) **0.5

    print("RMSE on training data: ", rmse_train)
    print("RMSE on test data: ", rmse_test)

    return model
```

```
[0]: model=perform_model(X_train2, y_train2, X_valid2, y_valid2, X_test2, y_test2)
```

```
RMSE on training data:  1.535566132453336
RMSE on test data:  1.8444138832320554
```

4.4 Part 3.4 CatBoost Feature Importance Discussion

```
[0]: feature_score = pd.DataFrame(list(zip(X2.dtypes.index, model.
    ↳get_feature_importance(Pool(X2, label=y2,
    ↳cat_features=categorical_features_indices))))),
    columns=['Feature', 'Score'])

feature_score = feature_score.sort_values(by='Score', ascending=False,
    ↳inplace=False, kind='quicksort', na_position='last')
```

```
[0]: plt.rcParams["figure.figsize"] = (12,7)
ax = feature_score.plot('Feature', 'Score', kind='bar')
ax.set_title("Catboost Feature Importance Ranking", fontsize = 14)
ax.set_xlabel('')
```

```

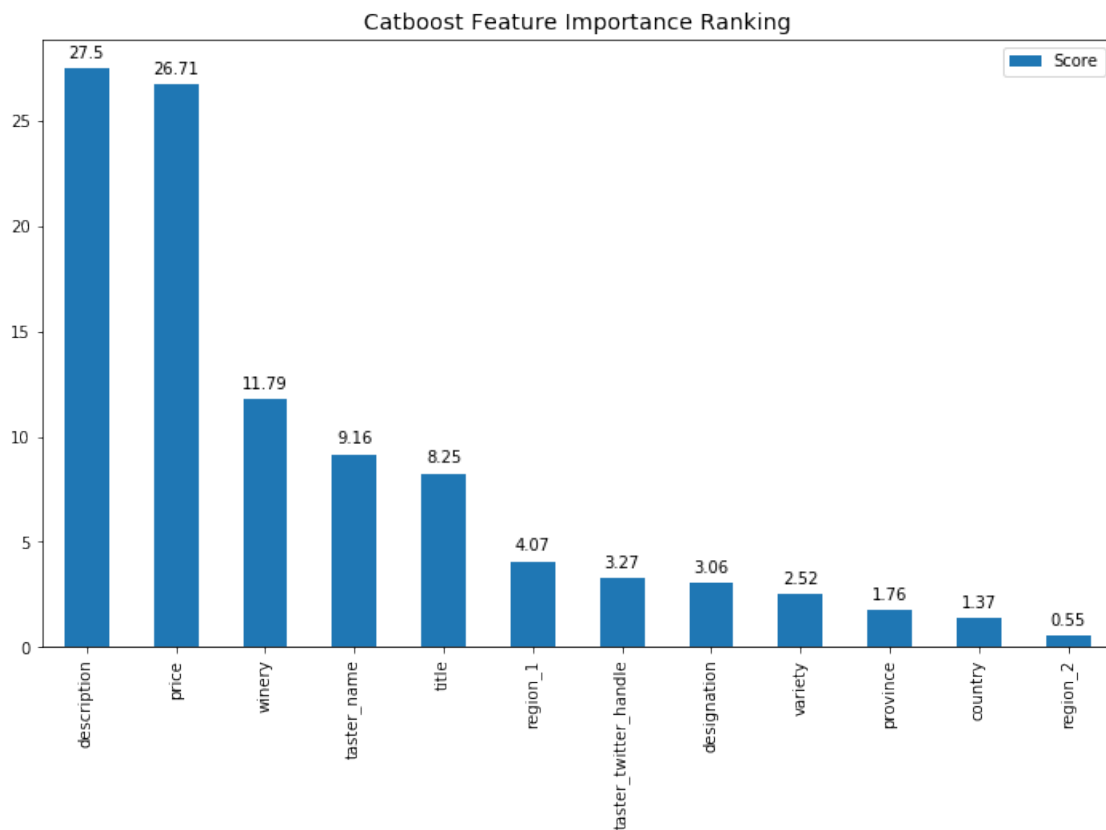
rects = ax.patches

labels = feature_score['Score'].round(2)

for rect, label in zip(rects, labels):
    height = rect.get_height()
    ax.text(rect.get_x() + rect.get_width()/2, height + 0.35, label,
            ↪ha='center', va='bottom')

plt.show()

```



As we can see from the feature importance plot, features 'region_2', 'country', 'province' and 'variety' are not important to CatBoost Model. Dropping those columns could decrease training time. Let's try to drop them and see how the score changes:

```

[0]: X2=X2.drop(columns=['country', 'province', 'region_2', 'variety'])

categorical_features_indices = [0,2,3,4,5,6]

X_train2, X_test2, y_train2, y_test2 = train_test_split(X2, y2, test_size=0.2,

```

```

random_state=42)

X_train2, X_valid2, y_train2, y_valid2 = train_test_split(X_train2, y_train2,
↳test_size=0.2,

random_state=52)

```

```
[0]: model=perform_model(X_train2, y_train2, X_valid2, y_valid2, X_test2, y_test2)
```

```

RMSE on training data: 1.5256643813745259
RMSE on test data: 1.8679501819819297

```

The RMSE of the model after I drop 4 columns which are not important doesn't change a lot. It's feasible to drop those columns to improve model efficiency.

5 Part 4: Small Recommendation System – Collaborative Filtering

This small recommendation system is performed by KNN, the similarity between each pair of customers is represented by cosine similarity. For a customer, I would like to recommend 5 wineries which are similar to the one the customer has reviewed. And those recommended wineries were tasted by 5 other customers who gave points higher than the average:

5.1 Part 4.1 Data Preprocessing

```

[0]: # filter the dataframe, only keep rows which satisfy the condition: 'pinot
↳noir', 'price less than 20 dollars', 'has a fruity taste' and 'points >
↳average'
def func(x):
    if 'fruity' in x:
        return True
    return False

X_re = pd.DataFrame(X[(wine_df['variety'] == 'Pinot Noir') & (wine_df['price']
↳< 20) & (wine_df['description'].apply(func)) & (wine_df['points'] >
↳wine_df['points'].mean())])
y_re = pd.DataFrame(y[(wine_df['variety'] == 'Pinot Noir') & (wine_df['price']
↳< 20) & (wine_df['description'].apply(func)) & (wine_df['points'] >
↳wine_df['points'].mean())]).reset_index(drop = True)
wine_re = wine_df[(wine_df['variety'] == 'Pinot Noir') & (wine_df['price'] <
↳20) & (wine_df['description'].apply(func)) & (wine_df['points'] >
↳wine_df['points'].mean())]
wine_re_train = pd.concat([X_re, y_re], axis = 1)

```

5.2 Part 4.2 Model Training and Recommendation

```
[0]: from sklearn.neighbors import NearestNeighbors
from scipy.sparse import csr_matrix
from sklearn.decomposition import TruncatedSVD

knn = NearestNeighbors(n_neighbors=10,algorithm= 'brute', metric= 'cosine')
model_knn = knn.fit(wine_re_train)
```

I randomly select a customer from raw dataset 'wine_df', and try to give him/her recommendation from the model.

```
[0]: import random
for i in range(5):
    query_index = np.random.choice(pd.concat([pd.DataFrame(X), pd.DataFrame(y)],
    ↪axis = 1).shape[0])
    distance, indice = model_knn.kneighbors(pd.concat([pd.DataFrame(X), pd.
    ↪DataFrame(y)], axis = 1).iloc[query_index,:].values.reshape(1,
    ↪-1),n_neighbors= 6)
    for i in range(0, len(distance.flatten())):
        if i == 0:
            print('Recmmendation for customer {0} in dataset row {1}:\n'.
    ↪format(wine_df.iloc[query_index, 8],pd.concat([pd.DataFrame(X), pd.
    ↪DataFrame(y)], axis = 1).index[query_index]))
        else:
            print('{0}: {1} with distance: {2}'.format(i, wine_re.iloc[indice.
    ↪flatten()[i], 12], distance.flatten()[i]))
```

Recmmendation for customer nan in dataset row 13883:

```
1: Dr. Nägler with distance: 0.011897025262433969
2: J. Lohr with distance: 0.015233030969037276
3: Willm with distance: 0.01562148796809526
4: Balletto with distance: 0.015658968179223076
5: Villa Wolf with distance: 0.015929275688086997
```

Recmmendation for customer Kerin O'Keefe in dataset row 51837:

```
1: Willm with distance: 0.00042388708821594623
2: J. Lohr with distance: 0.0007977300452143288
3: Three Brothers with distance: 0.0011844567231310554
4: Dr. Nägler with distance: 0.0012684173846814195
5: Murphy-Goode with distance: 0.0014919488310070061
```

Recmmendation for customer Lauren Buzzeo in dataset row 12932:

```
1: Willm with distance: 0.009167540030077137
2: Balletto with distance: 0.009586663853631494
```

3: J. Lohr with distance: 0.013522808583615209
4: Dr. Nägler with distance: 0.01485081544373379
5: Three Brothers with distance: 0.015207288103780137
Recommendation for customer Roger Voss in dataset row 83166:

1: Wakefield with distance: 0.04750312499019549
2: Willm with distance: 0.04791019183876277
3: Balletto with distance: 0.04878461196185424
4: J. Lohr with distance: 0.051494593165777425
5: Simonnet-Febvre with distance: 0.05534249118315393
Recommendation for customer Roger Voss in dataset row 55800:

1: Willm with distance: 0.006986447917567529
2: Balletto with distance: 0.00722844278658219
3: Esterházy with distance: 0.009294054246720362
4: P.J. Valckenberg with distance: 0.010290199278800038
5: J. Lohr with distance: 0.010525276043425125

This model could give a simple recommendation for a customer who has reviewed for any wines before. The model calculate the similarity between his/her review and other customers' reviews. This is not 'new-customer-friendly' because if a customer hasn't reviewed yet, the model won't work.